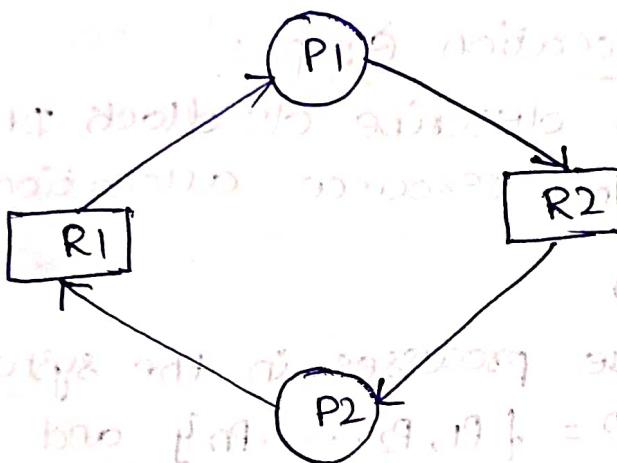


Deadlocks

- A deadlock is a situation where a set of processes are said to be in deadlock state if every process is holding some resources and is waiting for additional resources but those resources are acquired by other processes.



Deadlock characterization

Necessary conditions:

1. Mutual Exclusion
2. Hold and wait
3. NO preemption
4. circular wait.

① only one process can use only one resource at a time.

- That means all the resources are non-shareable by default.
- By default it is not possible to share resources among multiple processes as they lead in wrong results.

- ② A process is holding some resources and is waiting for additional resources.
- ③ No resource can be forcibly removed from a process holding it until it finishes its execution.
- ④ A set of processes are said to be in circular wait when P_1 holds R_1 and requests R_2 and process P_2 holds R_2 while it requests resource R_1 .

Resource - Allocation Graph:

- Is used to describe deadlock. It is also called system resource allocation graph.

$$G = (V, E)$$

All the active processes in the system are denoted by $P = \{P_1, P_2, \dots, P_n\}$ and set of resources denoted by $R = \{R_1, R_2, \dots, R_n\}$.

Process is denoted by circle

Resource is denoted by square

Dot within the square represents the no. of instances.

Request edge is an edge from process to resource

denoted by $P_i \rightarrow R_j$

Allocated edge or Assignment edge is an edge from resource to process denoted by $R_j \rightarrow P_i$

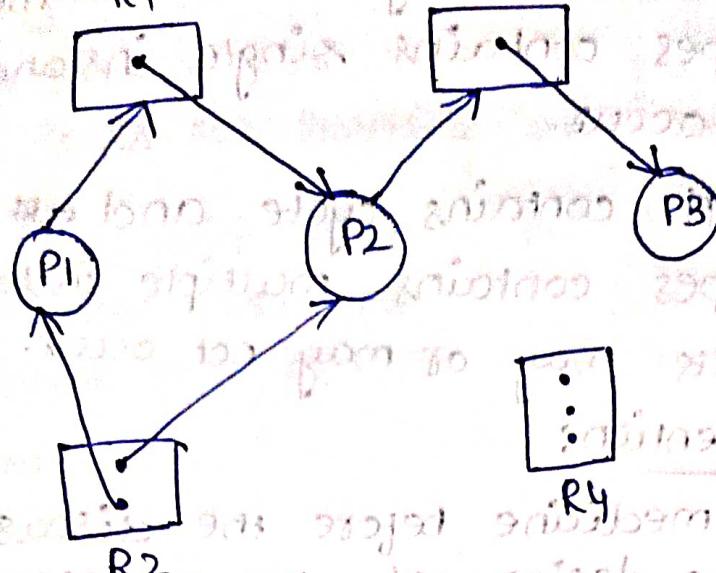
Ex:

$$P = \{P_1, P_2, P_3\}$$

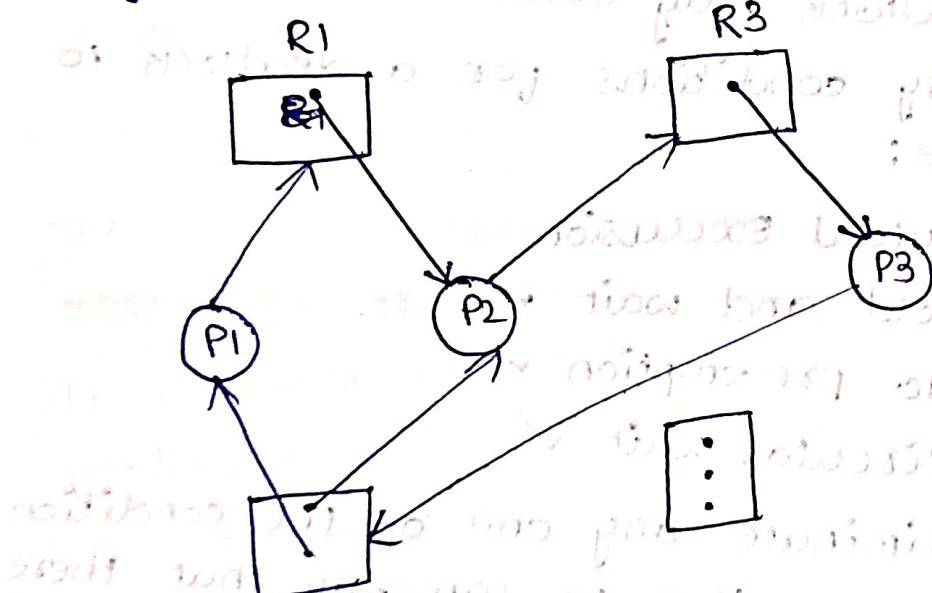
$$R = \{R_1, R_2, R_3, R_4\}$$

$$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_3, R_3 \rightarrow P_3\}$$

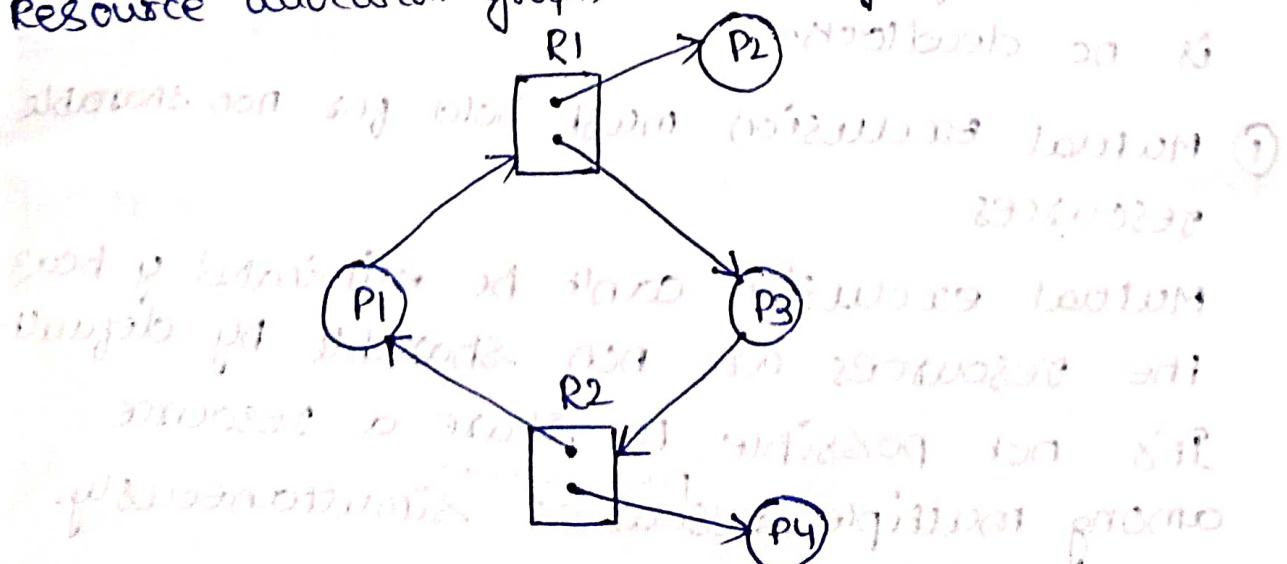
$$R_3 \rightarrow P_3\}$$



- If the graph contains no cycles, then no process in the system is deadlock.
- If the graph contains a cycle, then a deadlock may exist.



Resource allocation graph with a cycle but no deadlock:



- If the graph contains cycle and all the resources types contains single instance then deadlock ~~will~~ occur.
- If the graph contains cycle and ~~all~~ the resources types contains multiple instances then deadlock may or may not occur.

Deadlocks prevention:

- Taking the medicine before the disease.
- During the OS design only the corresponding OS manufacturers takes necessary precautions so that deadlock may not occur.
- If the 4 conditions occurs simultaneously in the system, then there is a possibility that deadlock may occur.
- Necessary conditions for a deadlock to occur are:
 1. Mutual Exclusion
 2. Hold and wait
 3. No preemption
 4. Circular wait.
- If we eliminate any one of the condition in the system, then we can say that there is no deadlock.

① Mutual Exclusion must hold for non sharable resources.

Mutual exclusion can't be eliminated bcz the resources are non sharable by default.

It's not possible to share a resource among multiple processes simultaneously.

- ② 2 Approaches:
1. A process can request for a resource only if it is not holding any other resources.
 2. The process has to put the request for all the resources.
 - once the process acquires all the resources then only the process has to start its execution.

Disadvantages:

1. Poor utilization of resources.

Ex: A process requires 5 resources and one among is printer which is required at the ending.

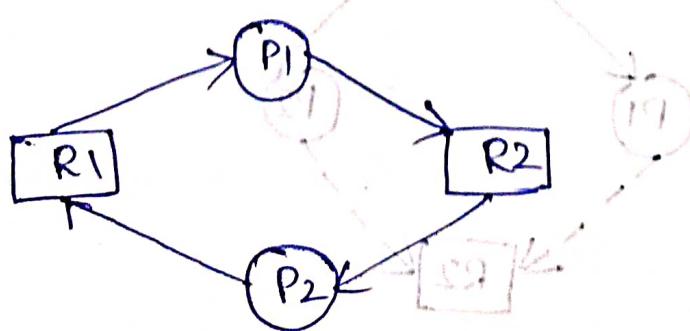
2. Starvation.

- ③ 1. $P_i \rightarrow R_j \rightarrow P_j$
- The OS will check whether P_i has any acquired resources or not. If P_i holds any other resources then OS forcibly releases those resources and allocates those resources to processes which need them. and after some time P_i may request for R_j .

2. $P_i \rightarrow R_j \rightarrow P_j \rightarrow R_k$
- $P_i \rightarrow R_j \quad P_j \rightarrow R_k$

Here, OS is forcibly releasing the resources.

④



- The corresponding process has to request for the process in incremental order
- $P_5 \rightarrow R_0, R_3, R_2$
- R_2, R_3, R_6

Deadlock avoidance:

1. Resource Allocation Graph Alg;
2. Bankers Alg. - multiple instances.

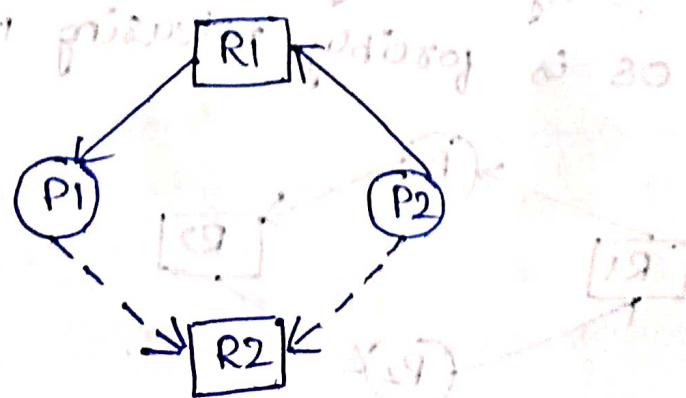
① If the resource type contains single instance, then we use Resource allocation graph alg. for deadlock avoidance.

- Normally request edge and assignment edge is used.
- In addition to the request edge and assignment edge, a new edge called claim edge is added.

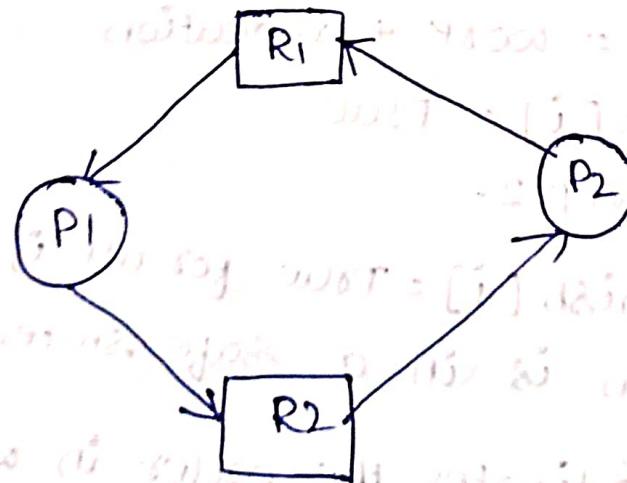
1. Claim edge - specifies that a process may request for the resource in the future.

- It is denoted by dotted arrow.

This indicates that process P_i may request for resource R_j in the future.



- when process P_i request resource R_j , then the claim edge is converted to a request edge.
 $P_i \rightarrow R_j$
- when the resource is free, then the OS allocates that resource to the process. Then the request edge is converted to assignment edge.
 $R_j \rightarrow P_i$
- once the process execution is over, then the process releases resources, so again that assignment edge is converted to claim edge.
- cycle detection alg. is used for detecting cycle in the graph.
- If no cycle exists, then the allocation of the resource will leave the system in a safe state.
- If a cycle is found, then the allocation puts the system in an unsafe state.
- safe state \Rightarrow No deadlock
- unsafe state - Deadlocks.



unsafe state in Resource Allocation Graph.

2. Banker's Algorithm: In order to avoid deadlock, there are two alg's:
1. Safety algorithm
 2. Resource - Request alg.
- Several data structures must be maintained to implement the bankers alg.
 - Let n be the no. of processes in the system and m be the no. of resource types. we need the foll. data structures.
- \rightarrow vector
 \rightarrow 1. Available - Indicates no. of available resources
- \rightarrow 2. Max - defines the max. demand of each process
- \rightarrow 3. Allocation - defines the no. of resources of each type currently allocated to each process
- \rightarrow 4. Need - indicates the remaining resource need of each process.

1. Safety Alg.:

Step1: work = Available

finish[i] = False for $i=0, 1, 2, 3, 4, \dots, n-1$

Step2: Find an i such that

a. $\text{Finish}[i] = \text{False}$

b. $\text{Need} \leq \text{work}$

Step3: work = work + allocation

$\text{Finish}[i] = \text{True}$

goto step 2.

Step4: If $\text{Finish}[i] = \text{True}$ for all i , then the system is in a safe state.

- Safe state indicates the order in which the processes are to be executed and OS will allocate resources to processes in that order only.

Example:

process	Allocation			Max			Available		
	R1	R2	R3	R1	R2	R3	R1	R2	R3
P0	0	1	0	7	5	3	3	3	2
P1	2	0	0	3	2	2			
P2	3	0	2	9	0	2			
P3	2	1	1	1	2	2			
P4	0	0	2	4	3	3			

SOL: The need is calculated as

$$\text{Need} = \text{Max} - \text{Allocation}$$

process	Need		
	R1	R2	R3
P0	7	4	3
P1			
P2	2	2	
P3	6	0	0
P4	4	3	1

work = Available

work = 3 3 2

Finish

False	False	False	False	False
0	1	2	3	4

P = 01 =

<09, 09, 09, 09>

P0

Need \leq work

7 4 3 \leq 3 3 2. False

So P0 must wait

P1

Need \leq work

1 2 2 \leq 3 3 2, True

work = work + Allocation

$$= 3 3 2 + 2 0 0$$

$$= 5 3 2 0 9$$

T	T	T	T	T
R	F	F	F	F
0	1	2	3	4

Finish[1] = True

P3

0 1 1 \leq 5 3 2, True

6 0 0 \leq 5 3 2 False

work = 5 3 2 + 2 1 1

P3 must wait

P4

Need \leq work

4 3 1 \leq 7 4 3, True

work = work + Allocation

$$= 7 4 3 + 0 2 0$$

$$= 7 4 5$$

Finish[4] = True.

P0

Need \leq work

7 4 3 \leq 7 4 5, True

work = work + Allocation

$$= 7 4 5 + 0 1 0$$

$$= 7 5 5$$

Finish[0] = True

P2

Need \leq work

1 2 2 \leq 7 5 5, True

work = 7 5 5 + 3 0 2

$$= 10 5$$

Finish[2] = True.

Safe sequence is $\langle P1, P3, P4, P0, P2 \rangle$

a. Resource - Request Alg:
 It is an extension to resource safety alg.
Step1: If $\text{Request}_i \leq \text{Need}_i$, go to Step 2
 Step 2.

Step2: If $\text{Request}_i \leq \text{Available}_i$, go to step 3.

Step3: $\text{Available}_i = \text{Available}_i - \text{Request}_i$

$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$

$\text{Need}_i = \text{Need}_i - \text{Request}_i$

If the resulting resource alloct. state is safe, the

Ex: Let P_1 request for $(1, 0, 2)$. Trans. is completed

1. $\text{Request} \leq \text{Need}$ and P_1 is allocated with its resources;

$$1 \ 0 \ 2 \leq 1 \ 2 \ 2 \text{ TRUE}$$

so goto step 2.

2. $\text{Request} \leq \text{Available}$

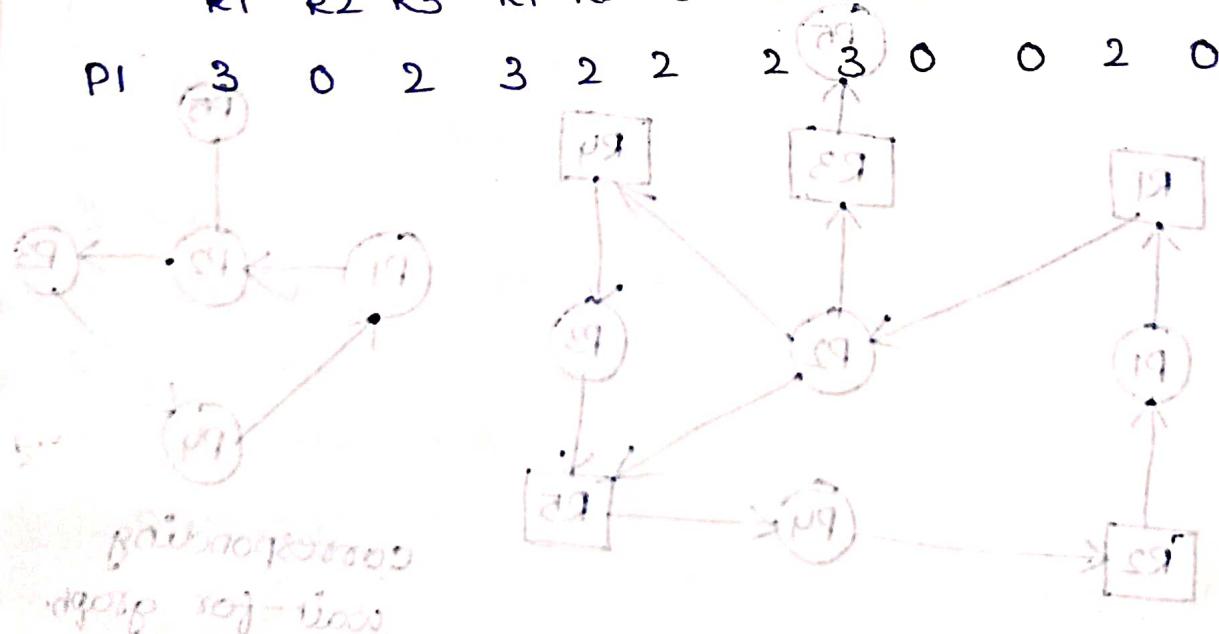
$$1 \ 0 \ 2 \leq 3 \ 3 \ 2 \text{ TRUE}$$

so goto step 3.

3. Allocation

	Max	Available	Need
R1	3	1	1
R2	0	2	2
R3	2	2	2

$$P_1 \quad 3 \quad 0 \quad 2 \quad 3 \quad 2 \quad 2 \quad 2 \quad 3 \quad 0 \quad 0 \quad 2 \quad 0$$



Eg: $P_1(1, 0, 2)$

Process	Allocation			Max			Available		
	R ₁	R ₂	R ₃	R ₁	R ₂	R ₃	R ₁	R ₂	R ₃
P ₀	0	1	0	7	5	3	2	3	0
P ₁	3	0	2	3	2	2			
P ₂	3	0	2	9	0	2			
P ₃	2	1	1	2	2	2			
P ₄	0	0	2	4	3	3			

Process	Need		
	R ₁	R ₂	R ₃
P ₀	7	4	3
P ₁	0	2	0
P ₂	8	0	0
P ₃	0	1	1
P ₄	4	3	1

i) Request_i ≤ Need_i

$$102 \leq 122 \text{ (True)}$$

Request_i ≤ Available

$$102 \leq 332 \text{ (True)}$$

$$\text{Available} = \text{Available} - \text{Request}_i$$

$$= 332 - 102$$

$$= 230$$

$$\text{Allocation}_i = \text{Allocation}_{i-1} + \text{Request}_i$$

$$= 200 + 102$$

$$= 302$$

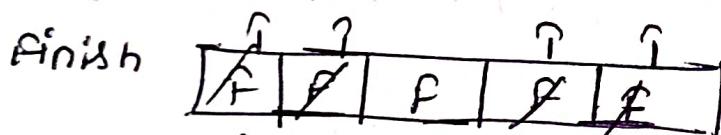
$$\text{Need}_i = \text{Need}_i - \text{Request}_i$$

$$= 122 - 102$$

$$= 020$$

work = Available

work = 230



$$P_0: \text{Need}^0 \leq \text{work}$$

$$743 \leq 230 \text{ (F)}$$

wait

$$P_1: \text{Need} \leq \text{work}$$

$$020 \leq 230 \text{ (T)}$$

$$\text{work} = \text{work} + \text{Allocation}$$

$$= 230 + 302$$

$$= 532$$

Finish[1] = true

$$P_2: \text{Need} \leq \text{work}$$

$$600 \leq 532$$

(False)

$$P_3: \text{Need} \leq \text{work}$$

$$011 \leq 532$$

True.

$$\text{work} = \text{work} + \text{Allocation}$$

$$= 532 + 211$$

$$= 743$$

Finish[3] = true

$$P_4: 431 \leq 743 \quad \boxed{1}$$

$$\text{work} = \text{work} + \text{Allocation}$$

$$= 743 + 002$$

$$= 745$$

Finish[4] = 1

$$P_0: 0.20 \leq 7.45 \text{ (r)} \quad P_2: 8.00 \leq 7.55 \text{ (r)}$$

$$\omega = \omega + A$$

$$= 7.45 + 0.10$$

$$= 7.55$$

$$\omega = \omega + A$$

$$= 7.55 + 3.02$$

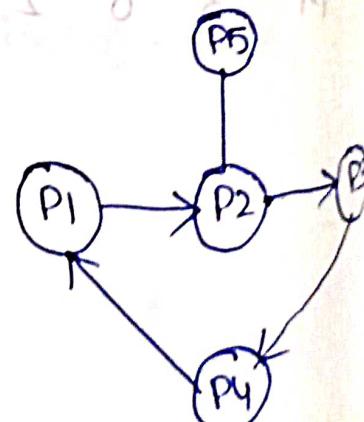
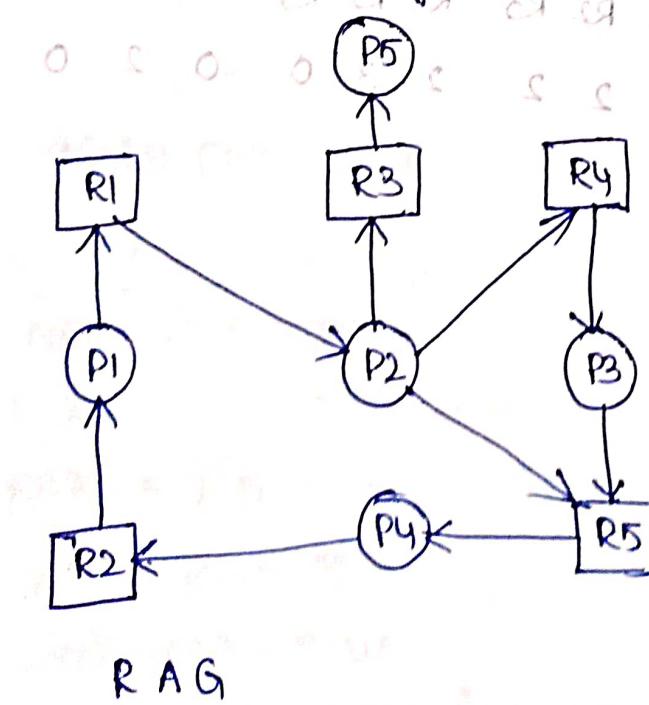
$$= 10.57$$

Finish[0] = true

∴ Sequence of safe : $P_1 \ P_3 \ P_4 \ P_0 \ P_2$.

Deadlock detection

- After allocating resources to the processes, the OS will check whether there is any deadlock or not.
- In order to detect deadlocks we have two approaches:
 - a. single instance of each resource type
 - a. Resource Allocation Graph
 - b. Several instances of a resource type
 - b. If all the resources have only a single instance, then deadlock detection alg. that uses a variant of RAG called wait-for graph.
 - c. wait-for graph is obtained from resource allocation graph.
 - d. Nodes of resource is removed and collapsing the appropriate edge i.e assignment edge and request edge.



corresponding
wait-for graph

- a. An edge from P_i to P_j in a wait-for graph implies that process P_i is waiting for process P_j to release a resource that P_i needs.
- b. An edge from P_i to P_j exists in a wait-for graph if and only if the corresponding resource-allocation graph contains two edges $P_i \rightarrow R_q$ and $R_q \rightarrow P_j$ for some resource R_q .
- If the wait for graph contains cycle, then there is a deadlock.
- To detect deadlock, the system needs to maintain the wait for graph and periodically to invoke an alg. that searches for a cycle in the graph.

2. several Instances:

Alg:

1. $\text{works} = \text{available}$
2. $\text{Finish}[i] = \text{False}$ if allocation $\neq 0$
 $\text{Finish}[i] = \text{True}$ if allocation $= 0$
3. Find an i such that

$\text{Finish}[i] = \text{False}$

Request \leqslant work

3. $\text{works} = \text{work} + \text{allocation}$

$\text{Finish}[i] = \text{True}$, goto step 2

4. If $\text{Finish}[i] = \text{False}$ for $i = 0, 1, 2, \dots, n-1$, then the system is in deadlock state

otherwise safe state.

	A	B	C
Allocation	7	2	6
Request	1	0	0
Available	0	0	0
P0	0	1	0
P1	2	0	2
P2	3	0	3
P3	2	1	0
P4	0	0	2

SOL:

$$WORKS = 0 \ 0 \ 0$$

Finish

F	F	F	F	F
0	1	2	3	4

P0

Request \leq WORKS

$$0 \ 0 \ 0 \leq 0 \ 0 \ 0, \text{ True}$$

$$WORKS = WORKS + ALLOCATION$$

$$= 0 \ 0 \ 0 + 0 \ 0 \ 0$$

$$= 0 \ 1 \ 0$$

Finish [0] = True

P3

Request \leq WORKS

$$2 \ 0 \ 2 \leq 3 \ 0 \ 3, \text{ True}$$

$$WORKS = WORKS + ALLOCATION$$

$$= 3 \ 0 \ 3 + 2 \ 1 \ 1$$

$$= 5 \ 2 \ 4$$

Finish [3] = True

P4

Request \leq WORKS

$$0 \ 0 \ 2 \leq 5 \ 2 \ 4, \text{ True}$$

$$WORKS = WORKS + ALLOCATION$$

$$= 5 \ 2 \ 4 + 0 \ 0 \ 2$$

P1

Request \leq WORKS

$$2 \ 0 \ 2 \leq 0 \ 1 \ 0, \text{ False}$$

So, P1 must wait

P2

Request \leq WORKS

$$0 \ 0 \ 0 \leq 0 \ 1 \ 0, \text{ False}$$

WORKS = WORKS + ALLOCATION

$$= 0 \ 1 \ 0 + 3 \ 0 \ 3$$

$$= 3 \ 1 \ 3, \text{ Finish}[2] = \text{True}$$

∴ No deadlock detected. Safe Sequence, $\langle P0, P2, P1, P4 \rangle$

P1

Request \leq WORKS

$$2 \ 0 \ 2 \leq 5 \ 2 \ 6$$

$$WORKS = 5 \ 2 \ 6 + 2 \ 0$$

$$= 7 \ 2 \ 6$$

$$Finish[1] = \text{True}$$

Recovery from deadlock

Two approaches

1. Process Termination

a. Abort all deadlocked processes

b. Abort one process at a time until deadlock cycle is eliminated.

2. Resource Pre-emption - 3 issues need to be considered

a. Selecting the victim

b. Roll back

c. Starvation

① a. 10 processes

3 P's are suffering from deadlock

P0 - 95% }
 P1 - 90% } To execute these processes it will require so much of
 P2 - 88% } CPU time, memory, I/O devices
 ... once we kill these processes then all those resources will be wasted. So

we can say that this is very expensive approach.

In this approach the resources will be wasted unnecessarily.

b. Kill one process at a time and by applying

deadlock detection technique in order to know deadlock cycle is eliminated or not.

- Too much overhead will be there on the

computer as we have to apply deadlock detection alg. after killing every process.

This is the dis. adv. of this approach.

- while aborting a process we must consider the foll. factors:

- * Lowest priority
- * Least amount of output produced so far
- * Least amount of processor time consumed so far
- * Least total resources allocated so far.

2. Releasing resources:

- a. which resources and which processes are to be pre-empted?
- b. Restarting all the processes from the beginning after pre-emption of resources or processes requires Rollback.
- c. There may be a possibility of starvation if a single process resources are released. In that case the process has to wait for longer period of time which leads to starvation.

By considering these three issues the resources have to be pre-empted from the process.

Methods for Handling Deadlocks:

1. Deadlock Prevention
2. Deadlock Avoidance
3. Deadlock Detection
4. Deadlock Recovery